JEDy: A Julia package for Evolutionary Dynamics

Nikolas M. Skoufis
Supervisor: Julián García

October 24th, 2014

Abstract

Evolutionary dynamics is the branch of biology concerned with studying how populations of invidivuals with inherited traits change over time. Performing simulations and calculations related to game theoretic views of evolutionary dynamics can be computationally intensive. In this project we have developed a library for performing common computational evolutionary dynamics calculations in the high performance scientific computing language Julia.

1 Introduction

Computational evolutionary dynamics studies the way that populations evolve when subject to inheritance of traits and mutation. The theoretical aspects of computational evolutionary dynamics are rooted in the mathematical discipline of game theory, however it is often useful to verify results with simulations. As the computations involved are generally repetitive and similar across various scenarios, it becomes useful to create packages of commonly used functions.

Packages for performing evolutionary dynamics calculations exist for languages such as Python [2] an Mathematica [4]. however these languages are notorious for being orders of magnitude slower than C or FORTRAN. However low level, high performance languages like C or FORTRAN do not offer the nice syntax and useful data strucures offered by higher level languages. In order to combine the performance of low level languages with the syntax of high level languages, we chose to write our package in Julia [1].

Julia is a relatively new language design for high performance scientific computing. Julia is syntactically inspired by Matlab and Python, but it remains fast. It is dynamically typed, uses multiple dispatch, is designed for parallelism and distributed computing, and benchmarks have shown that it can achieve near C speeds. Julia also has nice features like package management, access to iPython's notebook capabilities (for embedding rich text alongside code and figures), and the ability to call C and Python code. Julia is still under heavy development, and this created some problems during the development of this package.

2   Background

In order to understand what is involved in writing a package to perform computational evolutionary dynamics calculations, it will be useful to understand some of the game theory underpinnings of the calculations. While developing JEDy, we used the problem of the iterated prisoners dilemma in order to develop and test our package. Below we will present the problem and the mathematics behind simulating the system and solving for various quantities.

2.1   Iterated prisoner's dilemma

The prisoner's dilemma is a game involving two players. Each player is able to either cooperate or defect, and the four different combinations of these moves provides a different payoff to each player. If both players cooperate, they each receive payoff R. If both players defect, they each receive a lower payoff P. However if one player defects and the other cooperates, the cooperator receives the lowest payoff S, while the defector receives the highest payoff T. Therefore, we have payoffs such that $T < R < P < S$.

If we play the prisoner's dilemma multiple times, we have the ability to choose strategies. One strategy is to always cooperate (ALLC). Another possible strategy would be to allways defect (ALLD). The best strategy know is tit-for-tat (TFT) (or a minor variation of TFT) [3], where the player cooperates in the first game and then mimics the move that their opponent plays in the last game in subsequent games. Since TFT has some additional complexity over ALLC or ALLD, it is natural to introduce a complexity cost for this strategy. If we have a finite number of rounds m and a complexity cost c that reduces the payoff for TFT, we can construct a matrix which gives the payoff for each pairing of strategies.

$$
\begin{array}{c}
\begin{array}{ccc}
\text{ALLC} & \text{ALLD} & \text{TFT}
\end{array} \\
\begin{array}{c}
\text{ALLC} \\ \text{ALLD} \\ \text{TFT}
\end{array}
\left(
\begin{array}{ccc}
Rm & Sm & Rm \\
Tm & Pm & T+P(m-1) \\
Rm-c & S+P(m-1)-c & Rm-c
\end{array}
\right)
\end{array}
\tag{1}
$$

2.2   Reproduction, death and mutation: the Moran process

In studying the iterated prisoner's dilemma, we used the Moran process to model how populations of individuals evolve over time while playing the game. The Moran process can be classified as a birth-death process, where at each timestep the following steps occur in order:

1. An individual is chosen to reproduce. In reality this involves choosing a strategy to reproduce with probability proportional to their fitness (to be defined later).

2. The offspring of the reproducing individual may mutate to one of the other strategies with equal probability μ.

3. An individual is chosen to die. Each individual is chosen equiprobably.

The fitness function used in the Moran process is frequency dependant. Given a population with N individuals, with strategy populations x_1, x_2, x_3, and with payoff matrix A with elements a_{ij}, the fitness function can be defined as

$$f(i) = \frac{\sum_{j=0}^{3}(a_{ij}x_j) - a_{ii}}{N - 1}$$

for $i \in 1, 2, 3$. This is essentially the average payoff given that games are played against a random opponent. We subtract a_{ii} and divide by $N - 1$ because the player cannot play against themselves.

2.3   Intensity of selection

Another quantity of interest in computational evolutionary dynamics is the intensity of selection. The intensity of selection is a paramter which controls how much of a role fitness plays in selection. If we define p to be the payoff from the game and w to be the intensity of selection, we can transform the payoff values either linearly with

$$\pi(p) = 1 - w + wp, \quad w \in [0, 1]$$

or exponentially with

$$\pi(p) = e^{wp}, \quad w \in [1, \infty).$$

The benefit of transforming exponentially is that this allows us to deal with negative payoffs, where as the linear intensity of selection map cannot deal with negative payoffs.

As the intensity of selection approaches 0, all mapped payoffs approach 1. This creates a Moran process which is essentially a random walk between states, since at each timestep, each strategy is equally likely to be chosen to reproduce. As the intensity of selection approaches 1, the mapped payoffs approach their unmapped values, and as such, play a large role in selection.

Evolutionary dynamicists are often interested in the effect of intensity of selection on stationary distribution, which we will discuss in the next section.

References

[1] Jeff Bezanson, Stefan Karpinski, Viral B. Shah, and other contributors. The julia language. http://julialang.org, 2014.

[2] Julian Garcia. Evolutionary dynamics with python. https://github.com/juliangarcia/pyevodyn, 2014.

[3] Martin Nowak and Karl Sigmund. A strategy of win-stay, lose-shift that outperforms tit-for-tat in the prisoner's dilemma game. *Nature*, 364(1):56–58, 1993.

[4] Bill Sandholm, Emin Dokumaci, and Francisco Franchetti. Dynamo: Diagrams for evolutionary game dynamics. http://www.ssc.wisc.edu/~whs/dynamo/, 2014.